

Module	What Varies	Coverage	Headline Finding
Noise severity	Replay noise level (eps = 0.005-0.08)	MC / CP / LL; CBS, RV, DT	LunarLander mean BRA drop ranges 0.323-0.590; in MountainCar, RV and DT stay above 90% BRA through eps=0.08.
Rare-action support	Rare-action support and replay size	MC / LL quota resampling	MountainCar worst-R passes 0.70 at 5% support; LunarLander tops out at 0.341 even at 12%.
Cross-policy	Policy family with matched environments	CartPole + LunarLander; DQN vs PPO	CBS BRA changes only +0.006 on CartPole and +0.096 on LunarLander under PPO vs DQN.
Matching robustness	rho and lambda matching hyperparameters	CartPole + LunarLander	Matching choices matter: CartPole BRA spans 0.448-0.808; LunarLander mean F1 spans 0.412-0.523.
Bootstrap significance	Paired bootstrap confidence intervals and p-values	3 environments x 6 core claims	All 6/6 core claims are supported and significant in all 3 environments.
Repair-chain significance	DCM, MTC, and SSC repair stages	Merge-repair paired bootstrap (MC/CP/LL)	DCM->MTC is significant in 3/3 envs; DCM->SSC is significant in 3/3 envs; MTC->SSC is significant in 2/3.

All rows reuse existing experiment artifacts; the goal is to show the evaluation suite as a system rather than as isolated appendices.